

Question 1. Suppose you have a biased coin that when flipped takes on heads with unknown probability p and tails with probability $1 - p$. Show how to use this coin to construct a string of n independent bits such that each bit is equally likely to be a 0 or a 1. In other words, show how to use this coin to construct an algorithm that, when run, behaves like an unbiased coin. What is the expected running time of your algorithm as a function of p ?

Proof. We know that the consecutive sequence HT has the same probability as the consecutive sequence TH. So we would sample the biased coin 2 at a time, if we get HH or TT, we would throw the result out. If we get HT, then HEADS, otherwise if we get TH, then TAILS. Since $\frac{1}{2p(1-p)}$ is the number of pairs we expect to sample before we get either HT or TH. So the expected number of sampling the biased coin is $\frac{1}{2p(1-p)} \times 2 = \frac{1}{p(1-p)}$. \square

Question 2. [Inspired by Dave Moore] Imagine a procedure $\text{RANDOM}(a, b)$ that, when called, returns an integer between a and b inclusively and uniformly at random. That is, each integer in the range $[a, b]$ is equally likely to appear on a call to $\text{RANDOM}(a, b)$. Now, suppose you have a fair coin. Describe an implementation of $\text{RANDOM}(a, b)$ that is only allowed to flip this coin (i.e., it can't use any other source of randomness). What is the expected running time of your procedure, as a function of a and b ?

Proof. We will produce integers in the range $[0, b - a]$, and then map it into $[a, b]$. We will flip a coin $j = \lceil \log_2(b - a) \rceil + 1$ times, and let the result of each coin represent a binary bit (e.g. heads means 0 and tails means 1). We convert the resulting binary number into its decimal representation, d . We know that d is uniformly distributed in the interval $[0, 2^j - 1]$, so if $d > (b - a)$, then we rerun the algorithm until we can an integer in the range $[0, b - a]$. For each iteration we need the sample the coin j times, and we expect to have $\frac{2^j}{b-a}$ iterations. So the expected running time is $\frac{j2^j}{b-a} = \frac{(\lceil \log_2(b-a) \rceil + 1)2^{(\lceil \log_2(b-a) \rceil + 1)}}{b-a} \leq \frac{(\lceil \log_2(b-a) \rceil + 1)4(b-a)}{b-a} \approx O(\log(b - a))$. \square

Question 3. Consider a very simple online auction system that works as follows. There are n bidding agents; agent i has a bid b_i , which is a positive natural number. We will assume that all bids b_i are distinct from one another. The bidding agents appear in an order chosen uniformly at random, each proposes its bid b_i in turn, and at all times the system maintains a variable b^* equal to the highest bid seen so far (Initially b^* is set to 0).

What is the expected number of times that b^* is updated when this process is executed, as a function of the parameters in the problem?

Example. Suppose $b_1 = 20$, $b_2 = 25$, and $b_3 = 10$, and the bidders arrive in the order 1, 3, 2. Then b^* is updated for 1 and 2, but not for 3.

Proof. Let X_i be an indicator variable that we will have an update on the i th bid. We will only have an update on the i th bid if and only if the i th bid is the best amongst the first i bids. This happens with exactly probability $\frac{1}{i}$, since the best bid amongst the first i has equal probability in each of the i positions. So (expected number of updates) = $\sum_{i=1}^n X_i = \sum_{i=1}^n \frac{1}{i} \equiv O(\log n)$. \square

Question 4. Consider a county in which 100,000 people vote in an election. There are only two candidates on the ballot: a Democratic candidate (denoted D) and a Republican candidate (denoted R). As it happens, this county is heavily Democratic, so 80,000 people go to the polls with the intention of voting for D , and 20,000 go to the polls with the intention of voting for R .

However, the layout of the ballot is a little confusing, so each voter, independently and with probability $1/100$, votes for the wrong candidate – that is, the one that he or she didn't intend to vote for. (Remember that in this election, there are only two candidates on the ballot.)

Let X denote the random variable equal to the number of votes received by the Democratic candidate D , when the voting is conducted with this process of error. Determine the expected value of X , and give an explanation of your derivation of this value.

Proof. For sake of argument we will call people intending to vote for D democrats and people intending to vote for R republicans. $X = 20000 * (1/100) + 80000 * (99/100) = 79400$ because we expect 1% of republicans to wrongly vote for D and 99% of democrats to rightly vote for D . \square

Question 5. [extra credit] Suppose you have the same biased coin as in Question 1, however, this time you don't wish to construct an algorithm to produce an unbiased coin (i.e. a coin with probability $1/2$ of coming up heads) but rather another unbiased coin with some given probability q of coming up heads. Show how to use your coin to construct an algorithm that, when run, behaves like a biased coin with probability q of returning heads. You may assume that q is a rational number. Note: I don't know if this problem has a solution. Let's see.

Proof. From question (1) we know that we can simulate a fair coin using a biased coin with probability p of coming up heads. Now we simply need to show that we can use a fair coin to simulate a biased coin with probability q of returning heads. Since q is a rational number, let $q = \frac{a}{b}$, where a and b are coprime. Using the result of question (2) we know that we can use a fair coin to pick a random integer n uniformly distributed in the interval $[0, b]$. We can stipulate that if $n < a$, then HEADS, else TAILS, and we note that HEADS will occur with exactly probability $\frac{a}{b} = q$. The running time of this algorithm, combining the running time analysis of (1) and (2), is $\frac{O(\log b)}{p(1-p)}$. \square